JUSTIN P. DOLLMAN

Cambridge, UK | jpauldoll@gmail.com | **jdollman.github.io**

> *"User is highly engaged in the field of artificial intelligence, particularly AI policy, governance, and reasoning models. Their professional goal is to transition into the AI policy/AI governance space `Confidence=high`"*

> *"User has technical expertise in machine learning and deep learning, including interest in fine-tuning LLMs, reinforcement learning, and prompt engineering `Confidence=high`"*

— ChatGPT-4o[*]

[*] As the reader might know, LLMs now have "memories" for each user. This memory includes ChatGPT's `Model Set Context`, which collects knowledge that ChatGPT has gleaned from interactions with the client. For transparency, you can see my full "report" **here**.

PROFESSIONAL EXPERIENCE

Co-Director | **Cambridge AI Safety Hub (CAISH)**

October 2025 – Present | Cambridge, UK

- Co-directing operations and strategic initiatives for an AI safety field-building organization
- Director of MARS (Mentorship for Alignment Researchers): executed MARS 4.0 and organized MARS V
- Designed curricula for **Alignment Fellowship** (technical and governance cohorts)
- Co-organizes "How Do We Solve Alignment?" — weekly reading group of visiting researchers and Cambridge locals
- Founded the **Alignment Desk**

Lead Organizer | AI Safety Evals Group (BlueDot Impact)

June 2025 – October 2025

- As the sole paid organizer, directed a weekly forum for discussing cutting-edge research in AI Evaluations
- Secured guest speakers from leading AI safety organizations, including METR and MATS
- Revived a struggling group: attendance grew to its highest-ever levels under my leadership

Course Facilitator | ENAIS & AI Safety Hungary

July 2025 – September 2025

- Course: **Deep Dive: AI Policy 201**
- Led weekly discussion groups on intermediate AI governance topics

Course Facilitator | Center for AI Safety

June 2025 – September 2025

- Course: Introduction to AI Safety, Ethics, and Society (see: **here**)
- Lead weekly meetings for two cohorts and advised on participants' projects

NSF Researcher | Stony Brook University (Dept. of Technology and Society)

2020 – 2022

- Conducted bibliometric analysis and regression modeling on large datasets of NSF grant outcomes
- Cleaned, merged, and analyzed data using R to identify factors associated with successful broader societal impacts
- Contributed analyses and visualizations to grant reports and presentations

Graduate Teaching Instructor | SUNY Stony Brook (Dept. of Political Science)

2019 – 2025

- Instructor of record for Experiments in Political Science and Introduction to Statistics courses
- Effectively communicated complex statistical and programming concepts (R, causal inference) to diverse graduate and undergraduate audiences
- Developed and led workshop series: "Introduction to R for Social Sciences"

Wyzant Tutor | Remote

2021 – 2023

- Provided tutoring in statistics, linear algebra, calculus, research methods, and data visualization

EDUCATION

PhD Candidate, Political Science | Stony Brook University

2018 – Present (On Hold)

- Coursework focuses: machine learning and big data, causal inference, experimental design, American political behavior
- Dissertation research: urban-rural ideological divide, immigration attitudes, and subnational policy making
- Status: On hold to pursue full-time work in AI safety

BA, Psychology and Philosophy | University of Arkansas

2009 – 2014

- *Summa Cum Laude*
- Thesis: **The SNARC effect and the cognitive representation of a mental number line**
- Recipient of the J. William & Elizabeth W. Fulbright Scholarship (2014–2015)
  - Studied at Universität Karl-Franzens (Graz, Austria)
  - Specialized in philosophy of science

PROJECTS

Automated AI News Tracker & Visualization | *May 2025*

- Developed automated pipeline using GCP Cloud Functions and Scheduler to track AI news coverage
- Integrated with NYT Developer APIs to retrieve and analyze news data
- Implemented text classification using LLM API to identify AI-related content

R Package: callm | **Creator & Maintainer** | *2024 – Ongoing*

- Developed an R package providing a provider-agnostic interface for LLM API calls, simplifying workflows for integrating LLMs into research

Immigration Opinion and Policy Analysis | *2023 – 2024*

- Estimated state-level public opinion on immigration using MRP techniques on decades of Cooperative Election Study data
- Modeled the impact of opinion shifts on state policy adoption using time-series cross-sectional analysis
- Received the 2023 NYAAPOR/PANJAAPOR Young Public Opinion Star award for 'Immigration Opinion in the 50 States'

AI Dislocation and Economic Ideology | *Jan 2025*

- Designed and analyzed survey experiment linking economic ideology (status vs. equality concerns) to public perceptions of AI's economic impact

Debunk Bot Study Reanalysis | *Jan 2025*

- Conducted quantitative reanalysis of experimental data on AI chatbot effectiveness in countering misinformation

Broader Impacts and Researcher Productivity | *GTM2021 Conference*

- Presented analysis of factors associated with broader societal impacts of NSF-funded research

PUBLICATIONS

- **Technical Reviewer:** *Deep Learning for Biology* (O'Reilly, 2025) — co-authored by DeepMind researcher Natasha Latysheva
- Feldman, S., Mérola, V., & **Dollman, J.** (2021). The Psychology of Authoritarianism and Support for Illiberal Policies and Parties. *Routledge Handbook of Illiberalism.*
- **Dollman, J.** & Levine, W. H. (2016). The Mental Number Line Dominates Alternative, Explicit Coding of Number Magnitude. *Quarterly Journal of Experimental Psychology*, 69(3), 403–409.

SKILLS

Programming: `R` (Advanced) `Python` (Proficient) `SQL` (Proficient)

Data Science: Statistical Modeling, Machine Learning, Causal Inference, Time-Series Analysis, Survey Design & Analysis

Cloud/Tools: Google Cloud Platform (Functions, Scheduler), AWS (Lambda, Boto3), Docker, Kubernetes, FastAPI, Qualtrics, Git/GitHub

Languages: English (Native), Spanish (Fluent), German (Proficient Reading)

CERTIFICATIONS

⚠️ *LinkedIn-style Cheap Credential Flaunting*

AI Governance & Policy

- National Security Law and AI (Vista Institute) - *Jul 2025*
- AI Governance Course - *Apr 2025*
- EU AI Act Fundamentals - *Apr 2025*

AI, Machine Learning & LLMs

- Deep Reinforcement Learning - *Jan 2025*
- Agents Fundamentals - *Feb 2025*
- Developing Applications with LangChain - *Feb 2025*
- Reinforcement Learning in Python - *Mar 2025*
- Associate AI Engineer for Developers - *May 2025*

Data Engineering & Analysis

- Data Engineer in Python - *May 2025*
- Associate Data Engineer in SQL - *May 2025*
- SQL Fundamentals - *May 2025*

Python Development

- Associate Python Developer - *Apr 2025*
- Building Applications with OOP in Python - *Apr 2025*

Vista Institute | BlueDot Impact | Hugging Face | DataCamp